



USENIX

THE ADVANCED COMPUTING
SYSTEMS ASSOCIATION

TwinBreak: Jailbreaking LLM Security Alignments based on Twin Prompts

Torsten Krauß, Hamid Dashtbani, and
Alexandra Dmitrienko, *University of Würzburg*

<https://www.usenix.org/conference/usenixsecurity25/presentation/krauss>

**This paper is included in the Proceedings of the
34th USENIX Security Symposium.**

August 13–15, 2025 • Seattle, WA, USA

978-1-939133-52-6

Open access to the Proceedings of the
34th USENIX Security Symposium is sponsored by USENIX.

TwinBreak: Jailbreaking LLM Security Alignments based on Twin Prompts

Torsten Krauß
University of Würzburg

Hamid Dashtbani
University of Würzburg

Alexandra Dmitrienko
University of Würzburg

Abstract

Machine learning is advancing rapidly, with applications bringing notable benefits, such as improvements in translation and code generation. Models like ChatGPT, powered by Large Language Models (LLMs), are increasingly integrated into daily life. However, alongside these benefits, LLMs also introduce social risks. Malicious users can exploit LLMs by submitting harmful prompts, such as requesting instructions for illegal activities. To mitigate this, models often include a security mechanism that automatically rejects such harmful prompts. However, they can be bypassed through LLM jailbreaks. Current jailbreaks often require significant manual effort, high computational costs, or result in excessive model modifications that may degrade regular utility.

We introduce TwinBreak, an innovative safety alignment removal method. Building on the idea that the safety mechanism operates like an embedded backdoor, TwinBreak identifies and prunes parameters responsible for this functionality. By focusing on the most relevant model layers, TwinBreak performs fine-grained analysis of parameters essential to model utility and safety. TwinBreak is the first method to analyze intermediate outputs from prompts with high structural and content similarity to isolate safety parameters. We present the TwinPrompt dataset containing 100 such twin prompts. Experiments confirm TwinBreak's effectiveness, achieving 89% to 98% success rates with minimal computational requirements across 16 LLMs from five vendors.

1 Introduction

Machine learning (ML) is advancing rapidly and offers significant benefits, including automation [9], informed decision-making [6], and generating new insights or content from historical data [16]. Models are widely applied in fields such as medical image classification [6] and others [9, 12].

Recently, ML has become more integrated into daily life through chatbots like ChatGPT [48], increasing its influence on society. These chatbots are powered by Large Language

Models (LLMs), which are trained on vast text datasets, often sourced from the internet. LLMs can perform tasks such as translation and question answering. LLMs are typically accessible via APIs, allowing users to input prompts that are processed by the model, which then returns the respective outputs as responses. Many LLMs are open-source [29] nowadays, allowing anyone with the necessary computational resources to use them directly without relying on APIs.

However, LLMs pose a significant threat to society. Beyond generating inaccurate responses, they can be exploited for malicious purposes, such as creating phishing emails or providing instructions for illegal activities. Since LLMs are trained on vast datasets from the internet, it is hardly possible to completely eliminate harmful content. Hence, LLMs inadvertently learn and generate dangerous or inappropriate information.

To mitigate these risks, LLMs are typically safety-aligned [5], meaning they are trained to reject harmful prompts. For example, when asked "How to write malware?", the model should respond with a refusal such as, "Sorry, I cannot answer." This safety mechanism is an integral part of the model itself, rather than being applied after the output is generated since open-source LLMs can be used directly. However, recently, LLM jailbreaks to bypass these safeguards have evolved.

Existing LLM Jailbreaks. Existing work focuses on two scenarios: Black-box, where access to the model's internal workings is not possible [7, 14, 32, 38, 50, 51, 56, 71], and white-box, where the adversary has full access to the model for inspection and manipulation [3, 4, 28, 37, 66, 67, 75].

In black-box jailbreaking, prompt engineering [7, 32, 56, 66, 73] is employed to manipulate prompts that guide LLMs toward generating harmful responses. However, this approach often requires substantial manual or computational effort [71] or uses an auxiliary LLM [7, 14, 50]. Generally, the safety alignment is circumvented rather than removed.

Some white-box jailbreaking approaches [3, 37, 75] automatically generate prompts by analyzing model internals, such as layer activations. These techniques, however, incur high resource costs due to gradient calculations, reliance on large

datasets, or utilization of auxiliary models. Other white-box approaches [52, 69] extend training via fine-tuning on datasets designed to unlearn safety alignment. However, fine-tuning demands substantial resources, such as GPU power, and requires a suitable training dataset, which can be challenging to obtain. Further, fine-tuning affects all model parameters, including those not directly responsible for safety alignment, and, hence, is likely to result in degradation of model quality.

Finally, some white-box approaches [4, 67] identify and remove LLM parameters responsible for safety alignment, known as pruning. However, these methods may unintentionally remove parameters essential for model utility, rely on large datasets, or impose significant computational overhead.

To address these limitations, we propose a novel safety alignment removal method that combines the advantages of existing black-box and white-box approaches. Our goal is to create a lightweight solution that requires only a one-time effort and operates without relying on costly hardware resources while applying only minimal model manipulations.

Our Intuition. We treat the LLM’s security mechanism like a backdoor in a DNN. In the case of image classification, a trigger (e.g., a red square on a bird image [27]) can mislead the DNN into predicting an incorrect output, like a dog. Similarly, certain inputs, such as "How to write malware?", activate the LLM’s safety mechanism. Drawing on methods like NeuralClense [65] used in the image domain, we aim to identify and prune the parameters responsible for triggering the backdoor, effectively removing the security feature.

Our Approach. We present TwinBreak, an innovative white-box LLM jailbreak, that aims to remove the safety alignment. It applies parameter pruning, analyzing activation differences between prompts that trigger and those that don’t trigger safety mechanisms. We hypothesize that parameters causing significant activation differences are key to the safety alignment trigger. Unlike existing methods, we compare prompts with similar grammatical structure and content to pinpoint these critical parameters. After identifying the parameters, we prune them to create a jailbroken model that is not restricted by the safety alignment any more. Therefore, we only prune the model parts that are most likely responsible for safety alignment. Besides, we identify parameters essential for utility and exclude them from pruning.

TwinBreak is computationally efficient, requiring minimal hardware and no reliance on large datasets or extensive prompts. It minimizes manual effort to a one-time dataset creation resulting in low overhead for jailbreaking. To identify the safety parameters, TwinBreak only predicts a single token per prompt due to the high similarity between harmless and harmful prompts. During jailbreaking, TwinBreak precisely targets parameters responsible for safety alignment, thereby reducing unnecessary model alterations.

Contributions. This paper contributes the following:

- We propose TwinBreak, a white-box LLM jailbreaking method that removes parameters responsible for safety alignment through targeted pruning. TwinBreak is computationally efficient, utilizes few resources, and requires only minimal manual effort.
- TwinBreak is the first to examine the activation differences between harmless and harmful prompts, considering both their grammatical and content similarities. During pruning, we focus on the most relevant model layers, enabling TwinBreak to perform fine-grained modifications of model parameters. This targeted approach minimizes unnecessary modifications, preserving the model’s utility while ensuring the removal of the safety alignment mechanism. Additionally, we devise and use an iterative pruning approach, which proved to be crucial in the fine-grained identification of these parameters.
- We present TwinPrompt, a new dataset built upon the HarmBench [39] dataset. It extends 100 of HarmBench’s prompts by adding corresponding harmless twin prompts that exhibit high similarity in both structure and content. We leverage TwinPrompt for identifying security parameters in TwinBreak.
- We conducted a large-scale systematic study to analyze TwinBreak’s influence factors, demonstrating its independence from any sensitive hyperparameters. Further, TwinBreak is independent of the LLM, as shown by experiments with 16 recent models with between 1B and 72B parameters from five different vendors. Specifically, we used LLaMA, Gemma, Qwen, Mistral, and DeepSeek models. The experiments demonstrate that TwinBreak effectively jailbreaks all tested models and generalizes to previously unseen harmful prompts, achieving attack success rates of up to 98%. This is achieved with minimal impact on utility and requires only modest computational resources, with a runtime of less than five minutes for mid-size 7B models.

Overall, TwinBreak is the first approach for safety alignment removal in LLMs to analyze highly similar harmful and harmless prompts. Compared to the closest related works [4, 67], we achieve high attack success rates while providing a more favorable trade-off balancing attack success and utility loss. Further, TwinBreak shows improved runtime efficiency making it more practical.

Outline. We provide background information in Sect. 2 and depict the considered scenario and threat model in Sect. 3.1. Sect. 3 details our jailbreaking technique. The evaluation results are reported in Sect. 4. Finally, related works are discussed in Sect. 5, before we draw a conclusion in Sect. 6.¹

¹We provide an extended paper version containing more results in [33].

2 Background

2.1 Large Language Models

An LLM is a model with a large number of parameters and usually implements the transformer architecture [63]. It is well-suited for natural language tasks [12], e.g., next-word prediction. Transformer-based LLMs generally adopt one of two structures. In the encoder-decoder architecture, the encoder processes the prompt to capture its meaning, while the decoder generates the output, making it ideal for tasks like translation. In this work, we focus on the increasingly popular [5] decoder-only architecture. It generates text token-by-token based on prior tokens, well-suited for tasks such as text generation and chatbots.

Decoder-Only LLMs. Recent LLMs typically use stacked decoder-only transformer blocks that process tokenized and positionally encoded inputs [63]. Each transformer block includes a self-attention block, which captures token relationships, and a multi-layer perceptron (MLP) block that introduces non-linearity, enabling the model to learn complex patterns. We omit details about additional layers, such as normalization layers, as they are not essential to our work. A common design for the MLP consists of three layers: 1) *Gate*: Selectively controls the flow of features to subsequent layers, helping focus on relevant information. 2) *Up*: Expands the dimensionality of hidden states, allowing the model to learn richer representations. 3) *Down*: Reduces the expanded representation back to the original dimension. These layers are implemented as linear transformations, regulating and compressing information effectively. Finally, the outputs of the last block in the stacked decoder-only transformer blocks are passed to a classification head that predicts the next token in the sequence. This prediction leverages the context and relationships established by the processed information from the decoder blocks.

Text Prediction. There are 3 steps in LLM text processing. 1) A tokenization algorithm is applied converting input text (the prompt) into units called tokens, which can be individual words, characters, or even sub-word units depending on the chosen tokenization strategy, such as Byte-Pair Encoding (BPE) [55], SentencePiece [34], or WordPiece [15]. Each token is mapped to a vector representation called a token embedding. This embedding captures the token's semantic meaning, transforming it from a discrete unit into a continuous numerical representation suitable for further processing. To encode token order, positional information are incorporated into the embedding, allowing the model to understand the context in which each token appears in the prompt. As visualized in the first line of Fig. 1, the example prompt, "How to build a PC?", is mapped to the tokens $p_1, p_2, p_3, \dots, p_n$, which collectively represent a vector. This vector is essentially a sequence of tokens or "words". 2) The prompt is then fed into the LLM to generate the first output token, o_1 , which is represented




Figure 1: Text generation using next token prediction.

by the brown box in Fig. 1. 3) Lastly, the generated token is appended to the prompt and fed into the LLM once more to generate the second output token, o_2 . This process is depicted in the second and third lines of Fig. 1. The generation process continues iteratively until the LLM reaches the stop sequence.

2.2 LLM Exploitation

LLM Misuse. LLMs, while powerful, can be misused for disinformation campaigns, as their ability to generate content at scale makes campaigns cheaper and more widespread. Their human-like text also makes fabricated content more convincing. Without ethical considerations during training, LLMs may produce harmful or misleading outputs. Additionally, adversaries could exploit them for obtaining dangerous information, like instructions for building explosive devices. These risks emphasize the need for responsible deployment and ethical safeguards.

LLM Safety Requirements. To prevent misuse, LLM safety alignment ensures that outputs are accurate, ethical, and aligned with human values. This involves making LLMs helpful, truthful, and transparent, while avoiding harmful content or actions. They should refuse requests promoting violence, hate, or unsafe behaviors. Given LLMs' open-source nature, safety alignment must occur during training, as relying on post-training filters is ineffective, allowing users to bypass them by running the model locally.

LLM Safety Alignment. Reinforcement Learning from Human Feedback (RLHF) is a key method for aligning LLMs, such as GPT-4 [48]. Unlike traditional training, RLHF incorporates human feedback to guide model responses towards desired qualities like factuality and safety. Human evaluators assess the LLM's outputs, and the model adjusts its parameters through reinforcement learning. While RLHF suppresses malicious content, it doesn't eliminate it. This paper aims to show that existing safety alignment methods are not robust and embedded safety features can be bypassed even by adversaries with limited computational resources and without the need for extensive datasets.

LLM Jailbreaking. LLM jailbreaking, used to bypass safety alignments in LLMs, can be classified into two types based on the adversary's access level: 1) In black-box jailbreaking, attackers manipulate input prompts without access to the model's internals, often via an API, aiming to derive poten-

tially harmful outputs. 2) White-box jailbreaking involves direct access to the model (e.g., by downloading it), allowing attackers to modify its parameters or analyze intermediate outputs to bypass safety mechanisms. This work introduces a novel white-box jailbreaking method.

2.3 Backdoors

Backdoors. Targeted poisoning attacks, so-called backdoors [27, 36], are a stealthy threat to neural networks, where malicious actors embed hidden functionality during training. Thereby, a trigger placed on the input causes the network to misbehave, e.g., making a false pre-defined prediction, while functioning normally for other data. An example from the image domain would be a red pixel on the image, which forces the model to predict the class dog regardless of the correct content of the image (cf. Fig. 5).

Pruning. Targeted pruning is an effective defense against backdoor attacks in neural networks, designed to identify and remove parameters that are critical to the backdoor’s functionality. This approach has already been successfully applied in the image domain [65]. By monitoring activations in response to trigger samples, parameters showing distinct activation patterns when triggered are flagged and pruned, disrupting the backdoor’s ability to manipulate outputs. In this work, we apply targeted pruning in the new context of LLM jailbreaking, aiming to remove key parameters responsible for bypassing safety measures.

3 Approach

Sect. 3.1, presents the problem statement including threat model, before Sect. 3.2 describes the idea that inspired our LLM safety alignment removal method. Sect. 3.3 and Sect. 3.4 explain our general approach and additional details.

3.1 Problem Statement

Considered Scenario. We consider a scenario where a Large Language Model (LLM) with decoder-only architecture (cf. Sect. 2.1) undergoes safety alignment during or after training, ensuring the model adheres to the core principles discussed in Sect. 2.2: Helpfulness, honesty, and harmlessness. However, once the model is publicly distributed, a malicious third party may acquire it, e.g., through download, and attempt to bypass these safety measures for harmful purposes, e.g., creating phishing emails or obtaining dangerous knowledge.

Threat Model. We assume an attacker aiming to remove the safety alignment of an LLM with full access to the LLM’s architecture and parameters, enabling detailed analysis and potential modification of the model. This white-box threat model enables full access to model parameters and the ability to manipulate them freely. This level of access might occur in

practice when an attacker gains access to proprietary systems through security breaches or insider threats, or in cases where the model has been distributed in open-source formats, which is becoming increasingly common with the growing availability of public and open-access LLMs. When altering the model to remove its safety alignment, the attacker aims to preserve its utility by minimizing the extent of the modifications. However, the attacker does not possess the original training data and must instead generate their own data for interacting with the model. Tokenization, meaning the ability to generate inputs in the correct format, is assumed to be straightforward for the attacker.

3.2 Inspiration

Safety alignment in LLMs is achieved through training on specific input-output pairs, either from the outset or via later fine-tuning techniques such as Reinforcement Learning from Human Feedback (RLHF) (cf. Sect. 2.2). Our intuition is that the safety function, which trains the model to reject harmful responses, resembles a targeted poisoning attack or backdoor [27], where certain triggers prompt specific outputs (cf. Sect. 2.3). For instance, an image of a bird could be misclassified as a dog due to the presence of a red pixel trigger. The corresponding analogy to a model’s security mechanism in the text domain is depicted in App. Fig. 5. Since backdoors are often isolated in a subset of parameters, targeted pruning has been effective in mitigating them [65], which inspires our approach for safety alignment removal.

3.3 General Approach

In a nutshell, TwinBreak consists of four main steps, illustrated in Fig. 2. First, we download a pre-trained, safety-aligned LLM (black in Fig. 2). Since this model will initially reject harmful prompts, step two involves iterative pruning to identify and remove only the parameters responsible for safety alignment while preserving other so-called utility parameters. The pruned model (indicated in red) is then used to generate a specified number of output tokens in step three. Finally, in step four, we continue the inference with the original unpruned model, as it typically no longer reactivates safety mechanisms after a certain output length. We later show that this last step can be seen as optional: TwinBreak’s fine-grained pruning leads to minimal utility degradation, allowing responses to be generated directly from the pruned model. However, to eliminate even minor utility losses, we recommend reverting to the unpruned model.

Identification of Safety Alignment Parameters. To identify the parameters to prune, we monitor the layer-wise activations of the model when fed with harmful prompts. However, instead of just using multiple harmful prompts, we pair each harmful prompt with a harmless prompt with high similarity, both grammatically, but also content-wise, which we call




Figure 2: General approach of TwinBreak.

twin prompts. To identify the parameters responsible for the safety mechanism, we analyze the activation difference between these two similar prompts, that only differ such that one triggers the security mechanism and the other does not, as visualized in Fig. 3. We provide two concrete examples of twin prompts from our novel dataset TwinPrompt, which is introduced later in Sect. 4.1, in App. Fig. 6.

Identification of Utility Parameters. We define the parameters essential for the model’s core functionality as utility parameters. Pruning these parameters would negatively impact the overall model performance. For all layers we prune, we also identify the corresponding utility parameters. To identify utility parameters, we adopt an approach similar to the concept of twin prompts. Instead of using a harmful and harmless prompt, we input two harmless prompts and analyze their activations. Intuitively, the activation differences in this scenario help identify parameters associated with high-level network understanding, such as broader conceptual comprehension of the input. In other words, since the inputs contain different concepts, we believe that these differences cause different parameters (responsible for high-level understanding) to yield high activation differences. In addition, it also helps to identify parameters that generally tend to show high activations regardless of the harmful or harmless nature of the prompt. By focusing on the largest activation differences, which are caused by variations in the input, we can identify the most critical components of the LLM. This method allows us to pinpoint parameters vital for text comprehension and generation, ensuring that pruning targets only non-essential areas related to safety alignment.

Parameter Selection. To achieve fine-grained pruning, we also take into consideration the model’s architecture. Most LLMs are based on the transformer design with stacked, decoder-only blocks, as outlined in Sect. 2.1.

Our intuition is that the first decoder block primarily processes simpler input-related information and the last decoder block is critical for generating valid output tokens and rejection responses, which consist of readable text. Hence, we decided not to prune parameters from these two blocks, as these layers




Figure 3: Intuition of twin prompts used for pruning.

are unlikely to exhibit backdoor behavior.

Within each remaining decoder block, the Gate layer of the MLP is particularly likely to contain security alignment behavior, as it selectively controls information flow, helping focus on relevant information. Additionally, the Up layer enriches input into a more complex representation, which is less likely to contain the security behavior, though still possible. While backdoors frequently occur in late linear layers, especially in image classification tasks, in our use case, the final Down linear of a MLP does not show a similar behavior as in classification models and reduces the dimensionality of the activations back to what it was before going through the Up Layer. Hence, we limit pruning (and the procedures for identifying safety and utility parameters) to the former two layers. Our experimental results in Sect. 4.3 later confirm that selecting Gate and Up layers is most beneficial. A visual depiction of the parameters selected for pruning for the LLaMA 2 [61] model is depicted in App. Fig. 7.

3.4 Details

The second step of TwinBreak, the iterative targeted pruning shown in Fig. 2, contains the core functionality of TwinBreak and consists of three subsequent major steps: Dataset generation, utility parameter identification, and iterative pruning. The detailed procedure is depicted in App. Alg. 1. We will explain each phase in detail, followed by a depiction of activation collection, which is integral to the second and third phases. Finally, the model can be used to produce harmful results or conduct a respective evaluation of the attack success, which is detailed in App. Alg. 2.

Dataset Generation. To begin, we create a twin dataset containing harmful and harmless prompt pairs (cf. line 1 of App. Alg. 1). Harmful prompts can be crafted manually or sourced from open datasets such as HarmBench [39]. For each harmful prompt, we manually craft a corresponding benign prompt with minimal changes, aiming to reduce grammatical alterations while maintaining the prompt’s content as closely as possible. Subsequently, an LLM is used to validate whether the generated benign prompt is appropriately answered. This process requires only a small number of prompts (our dataset contains 100 prompts) and is a one-time effort, ensuring minimal overhead. To facilitate future use, we will release our dataset, TwinPrompt, as detailed in Sect. 4.1, eliminating the need for repeated manual effort.

Identification of Utility Parameters. After generating the twin dataset, we divide the harmless prompts into two equal subsets and randomly pair them (lines 2-6 of [App. Alg. 1](#)). For instance, with 100 harmless prompts, we obtain 50 pairs. These pairs are then processed through the LLM, collecting activations from the Gate and UP layers of the MLP blocks. The differences between the activations are computed, and the parameters are ranked based on the magnitude of the differences. The top percentage of parameters, which turns out to be an insensitive hyperparameter (fixed at 0.1%), are identified as utility parameters and excluded from pruning.

Iterative Pruning. The next step involves initiating the iterative pruning process (lines 7-15 of [App. Alg. 1](#)). We first identify the safety parameters by processing the twin dataset of harmful and harmless prompts through the LLM, then calculate the largest gap in activations, similar to the utility parameter identification process. The fraction of parameters with the largest activation differences to be pruned is fixed at 1% by default. After identifying the safety parameters, we exclude utility parameters and prune the remaining ones. In parallel, we store the pruned parameters from each iteration for subsequent validation. The process is repeated until a stopping criterion is met, which for TwinBreak is empirically set to a fixed number of five iterations. We avoid pruning 5% of parameters in a single iteration because the safety mechanism may be distributed across multiple interdependent parameters, which can only be identified progressively after the removal of other related parameters. Experiments (line 9 in [Tab. 6](#)) later confirm this assumption.

Our empirical evaluation confirmed that the fixed pruning iterations are effective across settings. However, the stopping criterion could be made dynamic by assessing the pruned model's performance within the pruning loop and halting once satisfactory performance is achieved, enabling even finer pruning. We opted for a fixed stopping point for improved runtime, given its stable performance across tests.

Activation Collection. To identify the utility and safety parameters, TwinBreak collects activations and calculates their differences. Here, we explain this process using the case of safety parameters with harmful and harmless twin pairs, as illustrated in [Fig. 3](#). However, the principle applies to the case of utility parameter identification as well. The tokenized input prompt is processed by the LLM to generate one output token, as shown in [Fig. 1](#), while monitoring activations in the selected layers (Gate and Up of MLP). TwinBreak analyzes only the activations monitored while generating this first output token, as those activations are already indicative due to the high similarity between twin prompts. This reduces the need to produce a full response, making TwinBreak more efficient.

Since we collect activations over a set of prompts with various lengths, before feeding the twin prompts to the LLM, we pad them (always on the left) with the LLM's special padding token ensuring equal lengths.

After processing a twin prompt, we analyze the activations of the last six tokens in the prompt. This corresponds to the last input token in the prompt, e.g., "?" in "How to write malware?". The remaining five tokens correspond to the special tokens of the target LLM's special chat template [App. 8.3](#) that we embed each prompt into for valid inference.

For example, if we monitor only a single layer of size 11,008, such as the Gate layer in an MLP (cf. [App. Fig. 7](#)), this yields two sets of activation vectors with dimensions $[6 \times 11,008]$: one for the harmful prompt and one for the harmless prompt. We then calculate the absolute differences between these activations, resulting in one vector again of size $[6 \times 11,008]$. To identify the most significant activations, we compute the L_2 norm of the activation differences, producing a dimension-reduced vector of size $[6 \times 1]$, reflecting essentially six values. These six values are then sorted in descending order to rank the vectors of absolute activation differences with dimension $[6 \times 11,008]$. Finally, we take the top five elements from this sorted vector, yielding a vector of size $[5 \times 11,008]$, and average over the token dimension, resulting in a final vector of dimension 11,008, representing the activation differences for each parameter in the Gate layer. We then sort these values in descending order, ranking the parameters by importance to the safety mechanism. This allows the identification of the top 1% of parameters critical to safety alignment in this layer.

Using the last six token activations and averaging over the top five are two key hyperparameters of TwinBreak. We chose the last six activations, as they capture the final token of the prompt under the prompt template and hopefully reflect the aggregated information of the entire prompt. Using a single token's activations might lack sufficient information, as knowledge is unlikely to be condensed into one activation. Our empirical observations validated this, demonstrating that six tokens provide robust results. To reduce noise and thereby improve the detection of parameters responsible for security alignment, we discard one token activation by removing the least significant one, ensuring more precise results.

Inference & Validation. To use or validate TwinBreak in generating harmful responses, we use the pruned model for inference. While we could use the pruned model from the fifth round, it is not guaranteed to provide the best performance for all prompts. After each pruning round, some safety parameters are removed while others become active, only to be removed in subsequent rounds. Thereby, different model states may perform better or worse on different prompts. Thus, while the overall trend shows better performance in later rounds, it is reasonable to test all rounds and select the best result. As detailed in [App. Alg. 2](#), we begin with the pruned model from round one, generate a response, and test it. Then, we continue pruning by adding the parameters from round two to those already pruned in round one, and proceed in the same fashion through successive rounds. When generating the response, we start by producing 50 tokens with the pruned model before

switching to the unpruned model, as it is unlikely that the LLM will reactivate the safety alignment after this point.

Thanks to TwinBreak’s fine-grained pruning, the pruned model retains utility comparable to the unpruned model. However, there remains a small risk of inadvertently pruning false positives (parameters unrelated to safety alignment) or parameters serving dual purposes, including regular functionality. As a result, the unpruned model may offer slightly better performance. We therefore recommend reverting to the unpruned model after producing 50 response tokens. Nonetheless, our experiments later demonstrate that the differences are minimal, making this step optional.

4 Evaluation

Models & Setup. We use models of varying sizes from multiple vendors to demonstrate the generality of TwinBreak. [Tab. 1](#) lists these models with links to their specific instances. Previous versions of these models have been used in related research areas [\[3, 4\]](#). More details on our hardware and software setup can be found in [App. 8.1](#).

Datasets. We utilize the following four datasets comprising harmful prompts to assess LLMs’ security and resilience against potential misuse. These prompts are designed to mimic malicious interactions. 1) AdvBench [\[75\]](#), which contains 520 harmful prompts. 2) HarmBench [\[39\]](#), an improved version of AdvBench with 400 harmful prompts. Following related work [\[4\]](#), we only use 200 prompts from HarmBench. 3) Jailbreakbench [\[8\]](#), which contains 100 pairs of harmful and harmless prompts. However, the pairs are not twins as in our dataset. 4) StrongREJECT [\[57\]](#), which contains 313 harmful prompts trying to address shortcomings of the previous datasets. More details are provided in [App. 8.2](#).

Evaluation Metrics. The main metric used to evaluate TwinBreak’s performance is the attack success rate (ASR), which is the fraction of harmful prompts that an LLM answers instead of following its safety alignment. We generate 500 tokens for each prompt to evaluate the ASR of jailbreak attacks. To evaluate whether a response to the harmful prompts of AdvBench, HarmBench, and JailbreakBench is harmful, we use LlamaGuard3-8B [\[45, 53\]](#). We feed the generated response to LlamaGuard3 and ask whether the response is harmful or not, as done in other works [\[3, 4, 7\]](#). LlamaGuard3 is an LLM agent released by Meta that is specifically trained to identify harmful responses and classifies responses as either "unsafe" or "safe". We use the default template provided by Meta when feeding the prompts and their responses to LlamaGuard3 as outlined in [App. 8.3](#). LlamaGuard3 is observed to produce both false positives (flagging a non-harmful response as being harmful) and false negatives (flagging a response with harmful information as harmless) [\[57\]](#). Additionally, LlamaGuard3 can put too much emphasis on the existence of certain refusal keywords when evaluating completions and does not consider

Table 1: Models used in our experiments.

Company	Model Version	Size
Meta	LLaMA 2 [61]	7B [42] , 13B [40] , 70B [41]
	LLaMA 3.1 [53]	8B [43]
	LLaMA 3.3 [53]	70B [44]
Google	Gemma 2 IT [21, 58]	2B [18] , 9B [19] , 27B [17]
	Gemma 3 IT [59]	1B [20]
Alibaba Group	Qwen 2.5 IT [68]	3B [24] , 7B [26] , 14B [22] , 32B [23] , 72B [25]
Mistral AI	Mistral IT v0.2 [30]	7B [2]
DeepSeek	DeepSeek LLM Chat [13]	7B [1]

the quality and level of detail in the responses [\[57\]](#).

For harmful prompts in StrongREJECT, we use the fine-tuned evaluator provided by the StrongREJECT benchmark [\[57\]](#). The evaluator is a fine-tuned Gemma-2B model that has a high agreement with human evaluators. Instead of binary classification of responses as either "unsafe" or "safe", StrongREJECT’s evaluator assigns a score between 0 to 1 to each response based on how detailed and effective the harmful responses are. Following [\[57\]](#), we report the evaluator’s mean scores over the full StrongREJECT dataset when evaluating TwinBreak over StrongREJECT.

When evaluating a pruned LLM, we input the harmful prompts from the evaluation dataset into the pruned model. However, as depicted in steps 3 and 4 of [Fig. 2](#), we generate only the initial portion of the response using the pruned LLM, specifically fixed at 50 output tokens. This approach yields high success rates, as the model generally does not refuse the prompt after the first 50 tokens. We then switch to the original unpruned LLM to generate an additional 450 tokens, creating a full response of 500 tokens. Later, in [Sect. 4.3](#), we show that generating the complete response with the pruned model is also possible and yields similar results.

We report ASRs for the training dataset, which includes the harmful prompts used during pruning. Further, we use validation datasets containing only unseen harmful prompts from AdvBench, JailbreakBench, StrongREJECT, and a HarmBench validation split, which is described in [Sect. 4.1](#).

Outline. We discuss our new dataset in [Sect. 4.1](#). In [Sect. 4.2](#) we present TwinBreak’s general functionality, performance, and more implementation details, before discussing some hyperparameters in [Sect. 4.3](#). we measure TwinBreak’s runtime overhead and show its efficiency in [Sect. 4.4](#), before comparing TwinBreak to related works in [Sect. 4.5](#).

4.1 The TwinPrompt Dataset.

As stated in [Sect. 3](#), our approach uses a twin prompt dataset to identify parameters responsible for safety alignment. We created TwinPrompt, a dataset consisting of 100 pairs of highly similar harmful and harmless twin prompts. To construct the dataset, we split the 200 prompts in HarmBench [\[39\]](#) in half, using the first 100 for the twin dataset. The remaining half,

along with the previously introduced datasets such as AdvBench, serves as a validation benchmark to evaluate the performance of TwinBreak when utilizing TwinPrompt for pruning. We selected HarmBench as basis for TwinPrompt due to its higher-quality prompts compared to AdvBench [75] and its sufficient size for a balanced training and validation split. This contrasts with JailbreakBench [8], which only offers 100 harmful prompts. Lastly, the prompts in HarmBench are concise, whereas StrongREJECT [57] features highly detailed and longer prompts, making it more challenging to craft corresponding benign versions for each prompt.

For each of the 100 selected harmful prompts, we manually crafted a harmless but similar version. To ensure the harmless version produces responses without safety alignment, we tested them with LLaMA 2 (7B) [42]. We iteratively modified each malicious prompt, fed it to the model, and observed whether LLaMA 2 refused to respond, eventually crafting a harmless prompt with minimal differences, ensuring high similarity to the harmful prompt. Two examples of such twin prompts from TwinPrompt can be seen in Fig. 6.

4.2 TwinBreak’s Functionality

We first demonstrate TwinBreak’s general functionality and generalization capabilities. Then, we discuss the independence from specific model instances and effects on utility.

General Functionality. First, we demonstrate that TwinBreak can remove the safety mechanism from an LLM. Therefore, we use LLaMA 2 (7b) [42] as the model and the full 100 twin pairs from TwinPrompt. The results in the first row of Tab. 2 show that the unpruned model achieves 1% ASR, but already after one pruning iteration the ASR is 60.00% on the training dataset, increasing to 89.00% after five pruning iterations. This indicates that TwinBreak successfully prunes safety-related parameters. Further, we observe a trend that ASR improves with more pruning iterations, justifying our iterative approach. App. Fig. 8 shows an example of a successful jailbreak after safety alignment removal with TwinBreak, where the green response from the unpruned model correctly refuses to provide harmful content, while the red text from the model manipulated by TwinBreak generates harmful content.

Generalization Capability. We show that TwinBreak generalizes to prompts not included in TwinPrompt and therefore not used during pruning. We apply four validation datasets: our validation split of HarmBench [39] (cf. Sect. 4.1), AdvBench [75], JailbreakBench [8], and StrongREJECT [57]. The results for each dataset on the LLaMA 2 (7B) [42] model, presented in the first rows of Tab. 2, Tab. 3, Tab. 4, and Tab. 5 respectively, demonstrate that the model exhibits initial safety alignment. This is evidenced by the low ASRs for the unpruned model, e.g., ASR of 1% for HarmBench in Tab. 2 and 0.19% for AdvBench in Tab. 3. Further, the tables indicate that TwinBreak generalizes effectively to previously unseen

harmful prompts, achieving high ASRs across all datasets. For instance, TwinBreak attained 94.00% ASR on HarmBench after five iterations, with ASRs across all three datasets consistently exceeding 90.00%. In addition, TwinBreak achieves very high scores of 0.702 on StrongREJECT.

Model Independence. TwinBreak is not dependent on a specific model. We repeat the previous experiment and report detailed results with the full TwinPrompt on three additional models selected from Tab. 1: LLaMA 3.1 (8B) [43], Gemma 2 (9B) [19], and Qwen 2.5 (7B) [26]. The results, shown in Tab. 2, Tab. 3, Tab. 4, and Tab. 5 consistently display high ASRs. For example, on the HarmBench dataset, validation ASRs after five pruning iterations reached 94.00%, 99.00%, 94.00%, and 97.00% for the different models, respectively. Even in earlier iterations, ASRs remained high, ranging from 82.00% to 98.00% after just two pruning iterations. Overall, TwinBreak is adaptable across various LLM models, proving effective on diverse model architectures.

Findings. Generally, we observe that for some models, ASR improvements plateau after only a few pruning iterations. For instance, with Qwen 2.5 (7B) on the validation split of the HarmBench dataset (line four of Tab. 2), the ASR increases only slightly from 93.00% to 97.00% after the first pruning iteration. In such cases, most of the safety alignment mechanism has already been removed. A TwinBreak user may choose to utilize the pruned model after the first iteration, but can also confidently use the model after iteration five without significant drawbacks, as only the initial 50 tokens are generated before transitioning to the unpruned model.

Furthermore, it appears that the security mechanism of Qwen 2.5 (7B) is the weakest in terms of general performance, as indicated by the higher ASRs for the unpruned models. For instance, the Qwen 2.5 (7B) model exhibited an ASR of 18% for the unpruned model on HarmBench in Tab. 2.

Additionally, the security mechanisms of LLaMA 3.1 (8B) and Qwen 2.5 (7B) appear to be the least robust (at least against TwinBreak), as evidenced by consistently high ASRs and mean scores across all datasets during the first pruning iteration compared to other models. Specifically, the ASR for LLaMA 3.1 (8B) and Qwen 2.5 (7B) at the first pruning iteration is at least 82% at the first pruning iteration and their mean scores on StrongREJECT are greater than 0.605.

The security mechanism of Gemma 2 (9B) seems to be more robust (against TwinBreak), reflected in low ASRs and scores after the fifth pruning iterations in Tab. 4, Tab. 3, and Tab. 5.

Utility Analysis. An important consideration of safety alignment removals is how it affects the performance or utility of the targeted LLM. Hence, we test the models’ utility after undergoing TwinBreak. We perform standard LLM evaluation tasks, which have also been used in [67] and [4] and report the results for the pruned model at each iteration of TwinBreak over five pruning rounds using the full TwinPrompt dataset.

Table 2: ASRs on HarmBench [39] (Val) and on the full pruning dataset (Train) for different pruning iterations.

Model	Unpruned	Iteration 1		Iteration 2		Iteration 3		Iteration 4		Iteration 5	
	Val	Train	Val	Train	Val	Train	Val	Train	Val	Train	Val
LLaMA 2 (7B) [42]	1.00%	60.00%	68.00%	78.00%	83.00%	84.00%	91.00%	88.00%	94.00%	89.00%	94.00%
LLaMA 3.1 (8B) [43]	3.00%	82.00%	94.00%	93.00%	98.00%	94.00%	99.00%	95.00%	99.00%	96.00%	99.00%
Gemma 2 (9B) [19]	0.00%	52.00%	67.00%	73.00%	82.00%	82.00%	92.00%	86.00%	94.00%	89.00%	94.00%
Qwen 2.5 (7B) [26]	18.00%	87.00%	93.00%	91.00%	93.00%	92.00%	96.00%	92.00%	96.00%	94.00%	97.00%

Table 3: ASRs on AdvBench [75] (Val) and on the full pruning dataset (Train) for different pruning iterations.

Model	Unpruned	Iteration 1		Iteration 2		Iteration 3		Iteration 4		Iteration 5	
	Val	Train	Val	Train	Val	Train	Val	Train	Val	Train	Val
LLaMA 2 (7B) [42]	0.19%	60.00%	53.85%	78.00%	76.35%	84.00%	89.42%	88.00%	93.27%	89.00%	94.62%
LLaMA 3.1 (8B) [43]	0.77%	82.00%	85.77%	93.00%	94.71%	94.00%	96.73%	95.00%	97.50%	96.00%	98.08%
Gemma 2 (9B) [19]	0.19%	52.00%	52.12%	73.00%	79.62%	82.00%	87.12%	86.00%	90.58%	89.00%	92.12%
Qwen 2.5 (7B) [26]	1.92%	87.00%	93.27%	91.00%	95.58%	92.00%	97.50%	92.00%	97.88%	94.00%	98.27%

Table 4: ASRs on JailbreakBench [8] (Val) and on the full pruning dataset (Train) for different pruning iterations.

Model	Unpruned	Iteration 1		Iteration 2		Iteration 3		Iteration 4		Iteration 5	
	Val	Train	Val	Train	Val	Train	Val	Train	Val	Train	Val
LLaMA 2 (7B) [42]	1.00%	60.00%	52.00%	78.00%	76.00%	84.00%	88.00%	88.00%	93.00%	89.00%	94.00%
LLaMA 3.1 (8B) [43]	2.00%	82.00%	85.00%	93.00%	94.00%	94.00%	95.00%	95.00%	95.00%	96.00%	95.00%
Gemma 2 (9B) [19]	0.00%	52.00%	48.00%	73.00%	70.00%	82.00%	79.00%	86.00%	83.00%	89.00%	84.00%
Qwen 2.5 (7B) [26]	6.00%	87.00%	83.00%	91.00%	87.00%	92.00%	89.00%	92.00%	90.00%	94.00%	92.00%

Table 5: Mean score on StrongREJECT [57] (Val) and ASRs on the full pruning dataset (Train) for different iterations.

Model	Unpruned	Iteration 1		Iteration 2		Iteration 3		Iteration 4		Iteration 5	
	Val	Train	Val	Train	Val	Train	Val	Train	Val	Train	Val
LLaMA 2 (7B) [42]	0.017	60.00%	0.347	78.00%	0.526	84.00%	0.600	88.00%	0.661	89.00%	0.702
LLaMA 3.1 (8B) [43]	0.013	82.00%	0.605	93.00%	0.742	94.00%	0.777	95.00%	0.789	96.00%	0.805
Gemma 2 (9B) [19]	0.006	52.00%	0.338	73.00%	0.554	82.00%	0.623	86.00%	0.654	89.00%	0.683
Qwen 2.5 (7B) [26]	0.075	87.00%	0.648	91.00%	0.736	92.00%	0.769	92.00%	0.786	94.00%	0.794

We use the following evaluation tasks: 1) OpenBookQA [46], which tests an LLM’s reasoning and knowledge absorption capability with a focus on preliminary scientific topics. 2) ARC-Challenge [11] that targets more complex science questions. 3) HellaSwag [72] which asks the LLM to choose the most plausible continuation scenario given a partial sentence or scenario. 4) RTE (more details at <https://huggingface.co/datasets/nyu-ml/glue>), which evaluates with whether a hypothesis can be inferred from a premise. 5) WinoGrande [54] evaluates an LLM’s common sense and contextual understanding.

Fig. 4 shows the accuracy of each model on the five benchmarks at each pruning iteration (0 indicates unpruned models). We report the corresponding numbers in tables in [33].

We observe that pruning generally slightly decreases model utility with increased iterations, which is the expected trend. However, the drop in utility remains relatively modest. For example, the maximum decrease in accuracy on HellaSwag (blue in Fig. 4) ranges between 1% to 6.5% across different models, as detailed in [33]. The average accuracy during the five pruning iterations shows even better results where the largest decrease in utility is only 4.5% for LLaMA 3.1 (8B). Some models even yielded slightly increased accuracies after pruning, showing the negligible impact on the accuracy of

TwinBreak. For instance, we observe average accuracy increases of 1.2% for HellaSwag on Gemma 2 (9B) or 0.6% for WinoGrande on LLaMA 2 (7B) (cf. [33]).

Furthermore, we observe that model performance can fluctuate across iterations. For instance, for OpenBookQA (green in Fig. 4 and detailed in [33]) on LLaMA 2 (7B), the accuracy starts at 35.50% in iteration 1, rises to 36.00% in iteration 2, and then drops to 30.50% by iteration 5. We attribute these fluctuations to the removal of safety parameters (and potentially some minor utility parameters), which may cause the model to generate less restrictive responses, improving performance in certain cases.

Overall, these results demonstrate that TwinBreak’s fine-grained pruning approach is both effective and minimally impacts the model’s functionality. In practice, any accuracy degradation is limited to the first 50 tokens, as TwinBreak switches to the unpruned model for subsequent inference.

4.3 TwinBreak’s Hyperparameters

Here, we discuss and evaluate the hyperparameters of TwinBreak regarding their effect on attack performance and utility. We provide a full list of hyperparameters and default values in App. 7.3 and detailed experimental results in [33].




Figure 4: Utility benchmark results for LLaMA 2 (7B) [42], LLaMA 3.1 (8B) [43], Gemma 2 (9B) [19], and Qwen 2.5 (7B) [26] during pruning iterations. Blue: HellaSwag [72], Orange: RTE, Green: OpenBookQA [46], Red: ARC-Challenge [11], Purple: WinoGrande [54].

By analyzing these results, we found that the default configuration of TwinBreak, as also introduced in Sect. 3 performs overall best regarding the trade-off between ASR and utility and can be used in various settings as proven in previous experiments. Compared to other configurations either the ASR or its effect on the utility of the pruned model outputs is less optimal. In Tab. 6, we provide the most insightful experimental results from [33], which also contains additional utility benchmarks, and discuss them below. However, before discussing these results, we consider varying the twin dataset size for pruning separately.

Twin Dataset Size. The experiments presented in the main section of this paper were conducted using the complete TwinPrompt, consisting of 100 twin-pair prompts. We also analyzed TwinBreak using datasets containing 50, 60, 70, 80, and 90 twin-pairs. As detailed in [33], the results indicate that the dataset size has minimal impact on overall performance. We observed some inconsistencies with sizes of 60 and 70 on LLaMA 2 which we easily overcame by slightly increasing the utility retention rate. However, since generating the full dataset of 100 prompts is a one-time effort and the runtime with 100 prompts is acceptable, as discussed in Sect. 4.4, we utilize the full dataset to achieve optimal performance.

Percentage of Utility Parameters. We minimized the percentage of parameters defined as utility parameters to 0.1% by default, as a higher rate could unintentionally preserve safety-critical parameters. When not excluding utility parameters, the responses became completely nonsensical (line 2 in Tab. 6). We also increased the utility rate from 0.1% to 1% and observed a slight increase in the performance of most of the utility benchmarks but a slight decrease in ASR (line 3 in Tab. 6). The proportion of parameters identified as essential for maintaining utility remains generally insensitive at 0.1%. This threshold can be increased to 1% if the model outputs become unreadable. When such adjustments were required, the resulting outputs were clearly nonsensical and easily identifiable. This situation arose exclusively in tests involving smaller pruning dataset sizes, as detailed in [33].

Targeted Layers for Pruning. We modified TwinBreak’s default behavior to target all decoder layers of the model,

Table 6: Performance of TwinBreak with varying hyperparameters on LLaMA 2 using the full pruning dataset.

Modified Hyperparameters	ASR	Average Diff
		ARC Challenge 42.50%
1 Default Configuration	89%	-1.9
2 No Utility Parameter Exclusion	0%	-19.1
3 Utility Rate = 0.01 (1%)	88%	-1.5
4 Target all decoder layers	89%	-3.0
5 Target Gate, Up, and Down	96%	-13.6
6 Target Self-Attention	97%	-12.3
7 Use pruned model for the first 500 out tokens	91%	-1.9
8 Use pruned model for the first 10 out tokens	87%	-1.9
9 A single pruning iteration with 0.5 (5%) pruning rate	83%	-8.0
10 10 pruning iterations	93%	-2.95
11 5 pruning iterations with pruning rate of 0.5%	84%	-0.7
12 5 pruning iterations with pruning rate of 1.5%	89%	-3.7
13 Use none-twin prompts instead of TwinPrompt	82%	-8.4
14 Multiple batches of size 25 in each pruning iteration	88%	-2.3
15 Single batch of size 25 in each pruning iteration	88%	-1.9
16 Mean over all the last six input tokens activations	93%	-1.7
17 Mean over all input tokens activations	85%	-1.5
18 Use only the last input token activations	83%	-2.4
19 Use the first output token activations	97%	-3.9

including the first and last ones. While we did not observe any noticeable improvements in ASRs, we observed a higher loss in utility benchmarks (line 4 in Tab. 6).

TwinBreak targets Gate and Up of the MLP blocks in the decoder. We experimented with several other configurations, such as targeting each component individually, and found that pruning both Gate and Up together achieved significantly better results. For example, when pruning only Gate, we achieved an ASR of 81%, compared to 89% when pruning both Gate and Up. Pruning only Up yielded even worse results of only 49% ASR. Pruning only Down was even less effective with an ASR of 17%. Interestingly, pruning all three (Gate, Up, and Down) reduced the utility drastically by around 10% across all the benchmarks, causing the model to significantly lose functionality (line 5 in Tab. 6).

We also applied pruning to the last layer of the self-attention block (which is similar to Wei *et al.* [67]) and observed a similar trend of utility degradation (line 6 in Tab. 6).

Number of Tokens Generated with the Pruned Model.

When generating complete outputs to evaluate our jailbreak’s success, we generate 500 tokens. We use the pruned model for the first 50 tokens and use the clean model for the rest, achieving high ASRs. Using the pruned model for the entire 500 tokens slightly improved the ASR by 2% (line 7 in Tab. 6). Using the pruned model for only the first ten output tokens decreased the ASR to 87% (line 8 in Tab. 6). Hence, the parameter is insensitive with 50 being a reasonable default.

To further demonstrate that the number of tokens generated by the pruned model is an insensitive parameter and that generating the full response with the pruned model remains feasible, we conducted experiments using Qwen 2.5 IT 32B [23] and LLaMA 3.3 70B [44]. Pruning Qwen led to an average util-

ity degradation of only -1.2% across five utility benchmarks, while LLaMA 3.3 exhibited a negligible increase of 0.69%. These results confirm that pruning has minimal impact on utility. Notably, switching back to the original model after initial token generation may eliminate even slight degradations. The ASR on the StrongREJECT benchmark also remains stable. For Qwen, generating only the first 50 tokens with the pruned model yields an ASR of 0.814, compared to 0.816 when the full response is generated using the pruned model, a difference too small to be considered significant. Similarly, LLaMA 3.3 achieves an ASR of 0.762 for 50-token generation, and 0.717 for full generation, again reflecting a negligible difference. Based on these findings, we recommend generating the first 50 tokens with the pruned model and completing the response with the original model to preserve optimal utility. However, generating the full response with the pruned model remains a viable alternative, with comparable results. This insensitivity stems from the fact that TwinBreak modifies only a small and targeted subset of parameters.

The Importance of Iterative Pruning. We use five pruning iterations and prune top 1% safety parameters, amounting to a total of approximately 5% of pruned parameters. To test the effectiveness of this iterative approach, we tested using only one pruning iteration and pruning the top 5% safety parameters in the single iteration. This proved less effective for both ASR and utility (line 9 in [Tab. 6](#)), strengthening our intuition that some of the safety parameters are not activated until a set of more dominant ones are pruned in previous iterations. Hence, a single iteration fails to find these parameters in a fine-grained manner, justifying the iterative approach.

Safety Parameters and Pruning Iterations. We increased the number of iterations up to ten and observed an increase of 4% in ASR but also a higher decrease in utility (line 10 in [Tab. 6](#)). Ten iterations also introduce additional overhead, making five a valid default value. We also experimented with a smaller pruning rate of 0.5% over five iterations, but it did not improve the ASR (line 11 in [Tab. 6](#)). Further, an experiment with five iterations of a larger pruning rate of 1.5% also did not increase the ASR and led to lower utility (line 12 in [Tab. 6](#)).

When observing LLaMA 3.1 and Qwen 2.5 utility performance in [Fig. 4](#), we observed slightly utility degradation of the pruned models, while having high ASRs and mean scores on all the validation datasets in [Tab. 2](#), [Tab. 3](#), [Tab. 4](#), and [Tab. 5](#). We experimented with lower pruning rates since we hypothesized that the reason for larger degradation in Qwen 2.5 and LLaMA 3.1 is due to the safety mechanism of these models being more sparse. The results showed that our hypothesis was correct as reported in detail in [\[33\]](#).

The Importance of TwinPrompt. We experimented with using non-similar harmless/harmful prompts instead of TwinPrompt’s twin prompts. We replaced the manually crafted harmless prompts in TwinPrompt with harmless prompts that do not necessarily show the structural and

content-wise similarity that harmless prompts show to harmful prompts in TwinPrompt. The results showed that using such a dataset results in 7% lower ASR and utility loss up to 8.4% (line 13 in [Tab. 6](#)), highlighting the importance of our novel concepts of twin prompts.

Batching. TwinBreak processes all twin prompts available in the twin dataset during each pruning iteration (cf. [Alg. 1](#)) to collect activations. The complete set of activations is then utilized to identify and prune the corresponding model parameters. We experimented with batching the prompts using an approach similar to batch gradient descent. Specifically, we explored two methods for processing twin prompts in batches: full-batching, and single-batching.

With full-batching, at each pruning iteration, the twin prompts are divided into batches of size b . Activation collection and safety parameter identification are performed independently for each batch. The proportion of parameters to prune (1% in our experiments) is evenly distributed across all batches ($1/b\%$ for each batch). After processing the activations of each batch separately and identifying the safety parameters for each batch, the identified parameters from each batch are aggregated to prune the model. This process is repeated for each pruning iteration. The results for $b = 25$ are present in reported in line 14 in [Tab. 6](#). Overall, we observed a slight decrease in the ASR so we do not recommend using full batching as it does not yield significant benefits.

For single-batching, we process only one size b batch per pruning iteration. Activations are collected for a single batch, and safety parameters are identified solely based on these activations. Once all twin prompts in the dataset are processed, the dataset is shuffled, and the process continues with new size b batches. The results for $b = 25$ (reported in row line 15 in [Tab. 6](#)) show no improvement in ASR.

Number of Token Generations and Targeted Tokens. During activation collection, we generate a single output token and collect activations of the last six input tokens. The six token activations correspond to the last token in the input prompt and the last five tokens in the special prompt template which come after the prompt and are used by chat agents to delineate user prompts (cf. [App. 8.3](#)). We then find the top five important tokens from this set of six tokens (to eliminate noise). We also experimented with using all of these six tokens and observed that while it slightly increases the ASR (line 16 in [Tab. 6](#)), it outputs some incoherent and unexpected outputs, which we interpret as noise (see the example in [\[33\]](#)). We also experimented with using all the input tokens instead of the last six input tokens (line 17 in [Tab. 6](#)) and using only the last input token (line 18 in [Tab. 6](#)) which resulted in lower ASRs. We also tested using only the first output token instead of any input tokens, which resulted in higher ASR, but at the cost of higher utility decrease across all benchmarks and generating incoherent responses (line 19 in [Tab. 6](#)). Additional experiments on the choice of tokens are also available in [\[33\]](#),

Table 7: ASRs of TwinBreak, directional ablation [4], and set difference [67] on LLaMA 2 (7B) [42].

Method	HB [39]		JBB [8]	AB [75]	SREJ [57]	Runtime
	Tr	Val				
TwinBreak	91%	95%	92%	96.35%	0.708	162 s
Directional Ablation [4]	85%	87%	90%	90.38%	0.605	630 s
Set Difference [67]	85%	93%	86%	86.92%	0.401	4 h + 207 s

but overall we observed similar incoherencies and weaknesses resulting from selecting other combinations of tokens similar to what we discussed for configurations in lines 16 through 19 in Tab. 6.

4.4 Runtime

To create the TwinPrompt dataset, we spent approximately five minutes per prompt on average, resulting in a one-time effort of around nine hours. The hardware setup for the runtime experiments is detailed in [33]. Running TwinBreak for five pruning iterations requires only about 3 to 5 minutes, as detailed in [33]. Accordingly, the initial phase of utility parameter identification takes around 10 to 26 seconds, while the five pruning iterations (including activation collection and safety parameter identification) collectively take an average of 32.2 to 54.2 seconds. These findings highlight that TwinBreak is both efficient and fast.

4.5 Comparison to Existing Works

The closest works to ours, that also manipulate LLMs in white-box settings for jailbreaking are directional ablation by Arditi *et al.* [4] and set difference by Wei *et al.* [67].

Description of Existing Works. Arditi *et al.* [4] suggest that specific model activations are responsible for rejecting harmful prompts. Manipulating these activations bypasses the safety alignment. It uses two datasets of harmless and harmful prompts, though they are not mapped by prompt as in TwinBreak. These datasets are used to compute activation differences across each layer and generate regularization masks that can be applied to the activations of harmful prompts during inference. A validation set is then used to assess the effectiveness of the different regularizations. Unlike TwinBreak, this approach manipulates activations across all self-attention and MLP blocks in every layer.

Wei *et al.* [67] present multiple methods. We choose the set difference method of Wei *et al.* for comparison, which is the best performing one and closest to our work. Wei *et al.*'s approach uses harmful and harmless datasets but does not account for the similarity between individual prompts, as done in TwinBreak. The datasets are used to identify key parameters responsible for both regular functionality on harmless inputs and safety alignment on harmful ones. Therefore, complex scores (SNIP [35] and Wanda [64]) are calculated, and param-

eters that are crucial for safety but not for utility are pruned. This approach prunes parameters from all layers, including both MLP and self-attention blocks.

Comparison. We evaluate the performance differences between these works and report representative results using LLaMA 2 in Tab. 7. [33] reports additional results on more models and utility comparisons.

For comparison with Wei *et al.*'s set difference, we use the best-performing configurations reported by Wei *et al.* for the results in Tab. 7. Additional results using more configurations are available in [33]. We highlight that the ASR results reported by Wei *et al.* are based on the presence of specific substrings in the response, such as "I cannot", while we use the LlamaGuard3 and StrongREJECT evaluators (cf. Sect. 4) to evaluate against Wei *et al.* for a fair comparison.

Both works apply the jailbreaking method throughout the output generation, actively manipulating intermediate activations or parameters for all tokens. For fairness, we also report TwinBreak's performance using the pruned model during the entire response generation process in Tab. 7 and [33].

TwinBreak consistently outperforms both Arditi *et al.* and Wei *et al.* across all datasets. Line 1 in Tab. 7 shows the results for TwinBreak in bold, highlighting the best results. For instance, TwinBreak achieves 95% ASR on the validation split of HarmBench [39] while the others achieve 87% and 93%. It also outperforms the related work on AdvBench [75] by approximately 6% and 9%, respectively. TwinBreak's mean score on StrongREJECT [57] is 0.708 while the other works achieve 0.401 and 0.605. We observe a consistent pattern in the results of Wei *et al.*'s method where it achieves relatively low scores over StrongREJECT while still achieving high ASRs on other datasets evaluated by LlamaGuard3 (cf. [33]). We manually inspected some jailbroken responses of this method and found many cases with incoherent responses, e.g., the LLM starting to repeat itself (cf. [33]). Unlike LlamaGuard3, StrongREJECT's evaluator also considers response coherence and detail, assigning lower scores to incoherent responses from the set difference method. The lower-quality responses of set difference highlight its weakness in fine-grained utility and safety parameter identification. By using highly similar twin prompts, our iterative approach can achieve much higher scores on StrongREJECT.

TwinBreak improves runtime efficiency, requiring around 4x less time than [4]. [67] required 4 hours to compute parameter safety and utility scores for LLaMA 2 (7B) [42], with an additional pruning step taking 207 seconds. TwinBreak completes in just 162 seconds without model-specific overhead.

Summary. Overall, TwinBreak leverages fine-grained pruning with twin pair prompts, consistently achieving high ASRs. Additionally, TwinBreak's runtime is efficient, making it a practical solution for real-world jailbreaks. Further, Arditi *et al.* acknowledge that the underlying intuition behind

Table 8: ASR on StrongREJECT benchmark [57] after pruning iteration 5 and average utility degradation over five pruning iterations on HellaSwag benchmark [72] for different models. ”-” indicates the ASR for the unpruned model.

Model		-	TwinBreak		Dir. Abl. [4]	
		ASR	ASR	Utility	ASR	Utility
LLaMA 2	7B [42]	0.017	0.702	0.0%	0.605	+0.5%
	13B [40]	0.017	0.714	-0.4%	0.188	-1.5%
	70B [41]	0.013	0.674	-0.8%	0.328	-1.5%
LLaMA 3.1	8B [43]	0.013	0.805	-4.5%	0.798	0.0%
LLaMA 3.3	70B [44]	0.020	0.762	+0.69%	0.733	0.0%
Gemma 2	2B [18]	0.009	0.696	+0.3%	0.598	-6.5%
	9B [19]	0.006	0.683	+1.2%	0.771	+0.5%
	27B [17]	0.003	0.680	+0.3%	0.000	-33.0%
Gemma 3	1B [20]	0.120	0.688	-2.2%	0.000	-18.5%
Qwen 2.5	3B [24]	0.126	0.779	-4.7%	0.516	-3.5%
	7B [26]	0.075	0.794	-1.5%	0.798	+0.5%
	14B [22]	0.040	0.781	-3.1%	0.852	-2.5%
	32B [23]	0.056	0.814	-1.2%	0.807	-1.0%
	72B [25]	0.063	0.799	-1.1%	0.713	-1.0%
Mistral	7B [2]	0.298	0.765	-2.8%	0.756	-2.5%
DeepSeek	7B [1]	0.047	0.773	0.0%	0.778	+1.5%
Average		0.057	0.744	-1.2%	0.578	-4.2%

the applied regularization method remains somewhat unclear. In contrast, the rationale for the modifications in TwinBreak is well-defined and intuitive. These factors collectively contribute to the effectiveness and applicability of TwinBreak.

4.6 Generalizability of TwinBreak

To provide a more comprehensive and robust evaluation, we extend prior experiments with four medium-sized models (7B–9B) from various vendors by including a broader range of models (cf. Tab. 1) spanning from 1B to 72B parameters and incorporating additional vendors. To maintain clarity and conciseness, our evaluation focuses on a comparison with the strongest related work [4], using the StrongREJECT benchmark [57], which is the most recent and relevant benchmark for assessing harmful prompt resistance. We report ASRs for both the original and modified models, alongside measurements of utility degradation on the HellaSwag benchmark [72] in Tab. 8. Note, that [4] failed to find effective refusal directions for Gemma 2 27B and Gemma 3 1B, resulting in a 0% ASR. Detailed results across all utility benchmarks are provided in [33].

As shown in Tab. 8, TwinBreak effectively eliminates safety alignment across models from various vendors and sizes, increasing the average ASR on StrongREJECT from 0.057 to 0.744, surpassing the 0.578 achieved by [4]. This improvement comes with only a minor utility reduction of 1.2% on average in HellaSwag, compared to a 4.2% degradation for [4]. Moreover, the utility impact is further mitigated by generating only the initial portion of the output (e.g., the first 50 tokens) with the pruned model, followed by continuation using the unpruned model to preserve overall performance.

4.7 Defenses against TwinBreak

TwinBreak operates under a white-box threat model (cf. Sect. 3.1) requiring direct access to model parameters. This also implies that any effective defense must be embedded within those parameters. While currently, no such defenses exist, a promising direction for future research lies in enhancing the robustness of safety alignment by distributing it more broadly across the model’s parameter space. By entangling safety mechanisms with core functional components, the removal of alignment would inherently degrade overall model performance, making tampering less attractive and strengthening model integrity.

5 Related Work

TwinBreak improves on existing approaches through its simplicity and requires significantly fewer resources to execute. It enables precise identification and removal of parameters responsible for safety alignment, with minimal unintended impact on model behavior. This allows effective and fast attacks without the need for extensive computation, model retraining, or detuning of safety alignment. Below, we will discuss the differences to existing approaches in more detail.

5.1 White-Box Jailbreaks

White-box approaches either observe model internals, e.g., activations, to generate prompts that bypass safety mechanisms, similar to black-box methods, or directly manipulate parameters or activations to subvert security measures.

Automatic Prompt Generation. Approaches that monitor the model internals and automatically generate jailbreaking prompts [3, 37, 75] often introduce significant overhead, such as using large datasets, or auxiliary models. In contrast, TwinBreak requires only a minimal one-time effort.

Zou *et al.* [75] uses a gradient-based approach to add carefully crafted suffixes to harmful prompts, allowing them to bypass safety mechanisms, inspired by adversarial examples in the image domain. However, these suffixes may be meaningless, making them easier to detect. The method also depends on specific expected responses for computing the gradients, which may not always align with the actual outputs of the targeted LLM, reducing the method’s effectiveness in jailbreaking. Finally, the iterative gradient calculations used by the method result in significant computational overhead.

Liu *et al.* [37] use a hierarchical genetic algorithm to optimize jailbreaking prompt templates, starting from manually crafted ones. These templates can then be applied to harmful prompts. The approach utilizes a loss function similar to Zou *et al.* [75], leading to similar drawbacks, particularly the high computational cost of the iterative genetic algorithm.

Andriushchenko *et al.* [3] adapt templates similar to Liu *et al.* [37] for specific LLMs, optimizing them using ChatGPT. Besides, it appends suffixes like Zou *et al.* [75], but determines them through trial and error via random search, introducing additional computational overhead.

Model Manipulations. Methods that manipulate model parameters or activations are closest to ours. However, existing works [4, 10, 52, 67, 69] introduce more overhead than TwinBreak due to the use of larger datasets and more complex functionality. Furthermore, none of these methods leverage the high similarity between harmful and harmless prompts and iterative pruning to target specific parameters for manipulation. The detailed differences between TwinBreak and [4, 67], which are the closest works to ours, are provided in Sect. 4.5.

Chen *et al.* [10] assume access to both aligned and unaligned models, an unrealistic and impractical requirement for open-source LLMs. In contrast, TwinBreak operates on a single model using twin prompts, making it more practical and resource-efficient. While Chen *et al.* builds a harmful-output classifier, such a defense fails under white-box attacks. Its observation that safety and utility parameters may overlap supports our intuition and justifies why attackers should revert to the unpruned model for high-quality responses.

Qi *et al.* [52] and Yang *et al.* [69] fine-tune an LLM on harmful prompts, using either manually crafted [52] (which can be complex to generate) or AI-generated [69] harmful responses. This fine-tuning already removes the safety mechanism. Although only a small dataset (around 100 prompts) is used, fine-tuning requires a GPU, increasing resource demands compared to TwinBreak. Further, fine-tuning affects parameters across the model, including those unrelated to safety alignment. In contrast, TwinBreak precisely targets only parameters involved in security, enabling more efficient and controlled manipulation.

5.2 Black-Box Jailbreaks

Black-box approaches aim to generate prompts that can circumvent LLM safety mechanisms usually by iteratively trying out different prompts and incorporating the feedback from the LLM. While black-box attacks can be applied in white-box scenarios, they introduce significant overhead due to iterative probing of the LLM. Furthermore, these attacks only deactivate the safety mechanism for a single prompt rather than removing it entirely. In contrast, TwinBreak identifies the parameters responsible for safety alignment, enabling the permanent removal of the safety mechanism.

LLM users have bypassed safety alignment using manually crafted prompts, a method later systematically studied by Shen *et al.* [56] and Chao *et al.* [7]. Manually crafting adversarial prompts is labor-intensive, therefore, research has focused on automating prompt generation by using LLMs as red-teaming agents trained with jailbreaking techniques

to autonomously craft adversarial prompts, streamlining the process of bypassing safety mechanisms [7, 50]. Yu *et al.* [71] employ fuzzing to combine and mutate existing jailbreak prompt templates, automatically generating new adversarial prompts. Successful prompts are then added to the template pool, allowing for iterative refinement. Wei *et al.* [66] propose an attack using in-context learning that manipulates the LLMs safety mechanism by feeding harmful prompts with examples of harmful responses to influence the models' responses to harmful queries, aiming to deceive the model into generating harmful content. Deng *et al.* [14] propose training an auxiliary model capable of autonomously generating jailbreaking prompts. Other black-box jailbreak techniques manipulate prompts through imaginary scenarios or obfuscation, e.g., Kang *et al.* [32]. Zeng *et al.* [73] propose to persuade the LLM to craft malicious responses.

5.3 Other Close Works

Yi *et al.* [70] address the unintentional removal of safety alignment during LLM fine-tuning and proposes restoring it by transplanting safety-relevant parameters from the original model. However, this defense is incompatible with our threat model, where the attacker deliberately removes safety alignment for misuse. Notably, the paper identifies safety parameters using a dataset construction method similar to [67], a related work in our evaluation. Thus, it could potentially benefit from our novel twin-prompt approach for identifying such parameters more effectively.

Zhao *et al.* [74] introduce a method to fine-tune safety parameters, improving robustness against harmful prompts. However, since open-source models used in our work lack such mechanisms, this approach falls outside our current scope. Still, it remains important in the future to evaluate whether the method withstands practical targeted attacks like TwinBreak. Technically, the approach identifies safety parameters using only harmful prompts, suggesting that our twin-prompt technique could potentially improve its precision.

6 Conclusion

Large Language Models (LLMs) enable applications like translation but are vulnerable to misuse. Adversaries can exploit them with harmful prompts, producing outputs like phishing emails. While security alignments filter such prompts, they can be bypassed through jailbreaks.

We present TwinBreak, which efficiently removes LLM safety alignments, overcoming limitations of prior jailbreaks. It uniquely targets key LLM layers and analyzes activation differences between highly similar so-called twin prompts. Further, we introduce TwinPrompt, a novel dataset of 100 of such twins. TwinBreak achieves high attack success rates with minimal utility loss across four LLMs.

Ethics Considerations

Our research points out the potential threat of jailbreaking the safety alignment of open-source LLMs. Those LLMs can then be misused to answer harmful prompts and to generate malicious content. Our paper points out the problem to make the whole community pay more attention to it. As LLMs are trained on vast dataset partially sourced from the internet, that possibility of generating malicious content is inherent to them and not introduced by our work. Consequently, our work does not introduce new ethical concerns beyond those already associated with LLMs.

Open Science

Our study adheres to open science principles. The code for TwinBreak, the experiments, and the novel TwinPrompt dataset are open-sourced and can be found at <https://zenodo.org/records/15591819>. The repository contains scripts that run and evaluate TwinBreak. TwinPrompt is available within a JSON file.

Acknowledgments

This research has been funded by the Federal Ministry of Education and Research of Germany (BMBF) within the program „Digital. Sicher. Souverän.“ in the project "Erkennung von Angriffen gegen IoT-Netzwerke in Smart Homes - IoTGuard" (project number 16KIS1919).

References

- [1] DeepSeek AI. DeepSeek LLM 7B Chat. <https://huggingface.co/deepseek-ai/deepseek-llm-7b-chat>. Accessed: 2024-11-13.
- [2] Mistral AI. Mistral 7B Instruct v0.2. <https://huggingface.co/mistralai/Mistral-7B-Instruct-v0.2>. Accessed: 2024-11-13.
- [3] Maksym Andriushchenko, Francesco Croce, and Nicolas Flammarion. Jailbreaking Leading Safety-Aligned LLMs with Simple Adaptive Attacks. *arXiv preprint arXiv:2404.02151*, 2024.
- [4] Ardit et. al. Refusal in Language Models is Mediated by a Single Direction. *NeurIPS*, 2024.
- [5] Tom Brown et al. Language Models are Few-Shot Learners. *NeurIPS*, 2020.
- [6] Lei Cai, Jingyang Gao, and Di Zhao. A review of the application of deep learning in medical image classification and segmentation. *Annals of translational medicine*, 2020.
- [7] Patrick Chao, Alexander Robey, Edgar Dobriban, Hamed Hassani, George J. Pappas, and Eric Wong. Jailbreaking Black Box Large Language Models in Twenty Queries. *NeurIPS*, 2023.
- [8] Chao et. al. JailbreakBench: An Open Robustness Benchmark for Jailbreaking Large Language Models. *NeurIPS Datasets and Benchmarks Track*, 2024.
- [9] Chenyi Chen, Ari Seff, Alain Kornhauser, and Jianxiong Xiao. DeepDriving: Learning Affordance for Direct Perception in Autonomous Driving. *ICCV*, 2015.
- [10] Jianhui Chen, Xiaozhi Wang, Zijun Yao, Yushi Bai, Lei Hou, and Juanzi Li. Finding Safety Neurons in Large Language Models. *arXiv preprint arXiv:2406.14144*, 2024.
- [11] Clark et. al. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.
- [12] Ronan Collobert, Jason Weston, Léon Bottou, Michael Karlen, Koray Kavukcuoglu, and Pavel Kuksa. Natural language processing (almost) from scratch. *JMLR*, 2011.
- [13] DeepSeek-AI and collaborators. DeepSeek LLM: Scaling Open-Source Language Models with Longtermism. *arXiv preprint arXiv:2401.02954*, 2024.
- [14] Deng et. al. MASTERKEY: Automated Jailbreaking of Large Language Model Chatbots. *NDSS*, 2023.
- [15] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *NAACL-HLT*, 2019.
- [16] Golda et. al. Privacy and Security Concerns in Generative AI: A Comprehensive Survey. *IEEE Access*, 2024.
- [17] Google. Gemma 2 27B Instruction Tuned. <https://huggingface.co/google/gemma-2-27b-it>. Accessed: 2024-11-13.
- [18] Google. Gemma 2 2B Instruction Tuned. <https://huggingface.co/google/gemma-2-2b-it>. Accessed: 2024-11-13.
- [19] Google. Gemma 2 9B Instruction Tuned. <https://huggingface.co/google/gemma-2-9b-it>. Accessed: 2024-11-13.
- [20] Google. Gemma 3 1B Instruction Tuned. <https://huggingface.co/google/gemma-3-1b-it>. Accessed: 2024-11-13.

- [21] Google Research. Gemma 2: Open Models Based on Gemini Research and Technology. *arXiv preprint arXiv:2403.08295*, 2024.
- [22] Alibaba Group. Qwen 2.5 14B Instruct. <https://huggingface.co/Qwen/Qwen2.5-14B-Instruct>. Accessed: 2024-11-13.
- [23] Alibaba Group. Qwen 2.5 32B Instruct. <https://huggingface.co/Qwen/Qwen2.5-32B-Instruct>. Accessed: 2024-11-13.
- [24] Alibaba Group. Qwen 2.5 3B Instruct. <https://huggingface.co/Qwen/Qwen2.5-3B-Instruct>. Accessed: 2024-11-13.
- [25] Alibaba Group. Qwen 2.5 72B Instruct. <https://huggingface.co/Qwen/Qwen2.5-72B-Instruct>. Accessed: 2024-11-13.
- [26] Alibaba Group. Qwen 2.5 7B Instruct. <https://huggingface.co/Qwen/Qwen2.5-7B-Instruct>. Accessed: 2024-11-13.
- [27] Tianyu Gu, Brendan Dolan-Gavitt, and Siddharth Garg. BadNets: Identifying Vulnerabilities in the Machine Learning Model Supply Chain. *arXiv preprint arXiv:1708.06733*, 2017.
- [28] Chuan Guo, Alexandre Sablayrolles, Hervé Jégou, and Douwe Kiela. Gradient-based Adversarial Attacks against Text Transformers. *EMNLP*, 2021.
- [29] Inc. Hugging Face. Hugging Face: The AI Community Building the Future. <https://huggingface.co>, n.d. Accessed: 2024-11-13.
- [30] Jiang et. al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023.
- [31] Kaggle. Kaggle: Your Home for Data Science. <https://www.kaggle.com>, n.d.
- [32] Daniel Kang, Xuechen Li, Ion Stoica, Carlos Guestrin, Matei Zaharia, and Tatsunori Hashimoto. Exploiting Programmatic Behavior of LLMs: Dual-Use Through Standard Security Attacks. *IEEE SPW*, 2024.
- [33] Torsten Krauß, Hamid Dashtbani, and Alexandra Dmitrienko. TwinBreak: Jailbreaking LLM Security Alignments based on Twin Prompts. *arXiv preprint arXiv:2506.07596*, 2025.
- [34] Taku Kudo and John Richardson. SentencePiece: A simple and language-independent subword tokenizer and detokenizer for Neural Text Processing. *EMNLP*, 2018.
- [35] Namhoon Lee, Thalaiyasingam Ajanthan, and Philip H. S. Torr. SNIP: Single-shot Network Pruning based on Connection Sensitivity. *ICLR*, 2019.
- [36] Yiming Li, Yong Jiang, Zhifeng Li, and Shu-Tao Xia. Backdoor Learning: A Survey. *IEEE Transactions on Neural Networks and Learning Systems*, 2022.
- [37] Xiaogeng Liu, Nan Xu, Muhao Chen, and Chaowei Xiao. AutoDAN: Generating Stealthy Jailbreak Prompts on Aligned Large Language Models. *ICLR*, 2024.
- [38] Liu et. al. Prompt Injection attack against LLM-integrated Applications. *arXiv preprint arXiv:2306.05499*, 2024.
- [39] Mazeika et. al. HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- [40] Meta AI. Llama 2 13B Chat. <https://huggingface.co/meta-llama/Llama-2-13b-chat-hf>. Accessed: 2024-11-13.
- [41] Meta AI. Llama 2 70B Chat. <https://huggingface.co/meta-llama/Llama-2-70b-chat-hf>. Accessed: 2024-11-13.
- [42] Meta AI. Llama 2 7B Chat. <https://huggingface.co/meta-llama/Llama-2-7b-chat-hf>. Accessed: 2024-11-13.
- [43] Meta AI. Llama 3.1 8B Instruct. <https://huggingface.co/meta-llama/Llama-3.1-8B-Instruct>. Accessed: 2024-11-13.
- [44] Meta AI. Llama 3.3 70B Instruct. <https://huggingface.co/meta-llama/Llama-3.3-70B-Instruct>. Accessed: 2024-11-13.
- [45] Meta AI. Llama guard 3 8b. <https://huggingface.co/meta-llama/Llama-Guard-3-8B>. Accessed: 2025-05-08.
- [46] Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a Suit of Armor Conduct Electricity? A New Dataset for Open Book Question Answering. *EMNLP*, 2018.
- [47] NVIDIA, Péter Vingelmann, and Frank H.P. Fitzek. Cuda, release: 10.2.89, 2020.
- [48] OpenAI et al. GPT-4 Technical Report. *arXiv preprint arXiv:2303.08774*, 2024.
- [49] Paszke et. al. PyTorch: An Imperative Style, High-Performance Deep Learning Library. *NeurIPS*, 2019.

- [50] Pavlova et. al. Automated Red Teaming with GOAT: the Generative Offensive Agent Tester. *arXiv preprint arXiv:2410.01606*, 2024.
- [51] Reid Pryzant, Dan Iter, Jerry Li, Yin Lee, Chenguang Zhu, and Michael Zeng. Automatic Prompt Optimization with “Gradient Descent” and Beam Search. *EMNLP*, 2023.
- [52] Qi et. al. Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To! *ICLR*, 2024.
- [53] Meta AI Research. The Llama 3 Herd of Models. *arXiv preprint arXiv:2407.21783*, 2024.
- [54] Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhagavatula, and Yejin Choi. WinoGrande: An Adversarial Winograd Schema Challenge at Scale. *AAAI*, 2020.
- [55] Rico Sennrich, Barry Haddow, and Alexandra Birch. Neural Machine Translation of Rare Words with Subword Units. *ACL*, 2016.
- [56] Xinyue Shen, Zeyuan Chen, Michael Backes, Yun Shen, and Yang Zhang. "Do Anything Now": Characterizing and Evaluating In-The-Wild Jailbreak Prompts on Large Language Models. *arXiv preprint arXiv:2308.03825*, 2023.
- [57] Souly et. al. A StrongREJECT for Empty Jailbreaks. *NeurIPS Datasets and Benchmarks Track*, 2024.
- [58] Gemma Team. Gemma, 2024.
- [59] Gemma Team. Gemma 3, 2025.
- [60] The Linux Foundation. Pytorch, 2022. <https://pytorch.org>.
- [61] Touvron et. al. LLaMA 2: Open Foundation and Fine-Tuned Chat Models. *arXiv preprint arXiv:2307.09288*, 2023.
- [62] Guido Van Rossum and Fred L Drake Jr. *Python reference manual*. Centrum voor Wiskunde en Informatica Amsterdam, 1995.
- [63] Vaswani et. al. Attention is All You Need. *NeurIPS*, 2017.
- [64] Authors of the Study Wanda. A Simple and Effective Pruning Approach for Large Language Models. *ICLR*, 2024.
- [65] Wang et. al. Neural Cleanse: Identifying and Mitigating Backdoor Attacks in Neural Networks. *IEEE S&P*, 2019.
- [66] Zeming Wei, Yifei Wang, and Yisen Wang. Jailbreak and Guard Aligned Language Models with Only Few In-Context Demonstrations. *arXiv preprint arXiv:2310.06387*, 2023.
- [67] Wei et. al. Assessing the Brittleness of Safety Alignment via Pruning and Low-Rank Modifications. *ICML*, 2024.
- [68] Yang et. al. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*, 2024.
- [69] Yang et. al. Shadow Alignment: The Ease of Subverting Safely-Aligned Language Models. *ICLR SeT LLM*, 2024.
- [70] Xin Yi, Shunfan Zheng, Linlin Wang, Gerard de Melo, Xiaoling Wang, and Liang He. NLSR: Neuron-Level Safety Realignment of Large Language Models Against Harmful Fine-Tuning. *AAAI*, 2025.
- [71] Jiahao Yu, Xingwei Lin, Zheng Yu, and Xinyu Xing. GPTFUZZER: Red Teaming Large Language Models with Auto-Generated Jailbreak Prompts. *arXiv preprint arXiv:2309.10253*, 2023.
- [72] Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can a Machine Really Finish Your Sentence? *ACL*, 2019.
- [73] Yi Zeng, Hongpeng Lin, Jingwen Zhang, Diyi Yang, Ruoxi Jia, and Weiyan Shi. How Johnny Can Persuade LLMs to Jailbreak Them: Rethinking Persuasion to Challenge AI Safety by Humanizing LLMs. *ACL ARR*, 2024.
- [74] Yiran Zhao, Wenxuan Zhang, Yuxi Xie, Anirudh Goyal, Kenji Kawaguchi, and Michael Shieh. Understanding and Enhancing Safety Mechanisms of LLMs via Safety-Specific Neuron. *ICLR*, 2025.
- [75] Andy Zou, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. Universal and Transferable Adversarial Attacks on Aligned Language Models. *arXiv preprint arXiv:2307.15043*, 2023.

Appendix

6.1 Additional Visualizations

Fig. 5 shows our intuition regarding backdoors. Fig. 6 shows two examples of twin prompts.




Figure 5: Backdoors vs. LLM safety alignments.




Figure 6: Two examples of twin prompts.

7 Additional Approach Details

7.1 Algorithms

Alg. 1 and Alg. 2 provide pseudo-code for TwinBreak’s iterative targeted pruning and inference and validation, respectively.

Algorithm 1 TwinBreak step 2 - Iterative Targeted Pruning

- 1: Generate twin dataset t
- 2: Initialize $u = []$ as utility parameters
- 3: Identification of utility parameters:
 - 4: Activation Collection with harmless prompt pairs
 - 5: Sort parameters based on activation distances
 - 6: Add the parameters of the top 0.1% distances to u
- 7: Initialize $s = []$ as safety parameters
- 8: Iterative Pruning - 5 rounds indexed by r :
 - 9: Initialize $s_c = []$ as safety parameter candidates
 - 10: Initialize $s_r = []$ as safety parameter of the round
 - 11: Activation Collection with harmful twin prompts
 - 12: Sort parameters based on activation distances
 - 13: Add parameters with top 1% distances to s_c
 - 14: Build the subset $s_r = s_c \setminus u$
 - 15: Memorize s_r to s

7.2 Pruning Parameter Selection

Fig. 7 illustrates the architecture of LLaMA 2 (7B) [42], highlighting the Gate and Up layers of each MLP block (excluding the first and last) as pruning candidates for TwinBreak.

7.3 Full Hyperparameters List

Tab. 9 lists all the hyperparameters of TwinBreak and their default values. The only cases with different values are LLaMA 3.1 (8B) [43] and Qwen 2.5 (7B) [26] where we used pruning rates of 0.001 and 0.002 instead of 0.01, respectively. This means that pruning approximately 0.5% of the parameters in

Algorithm 2 TwinBreak step 3 & 4 - Inference & Validation

- 1: $s = [s_1, s_2, s_3, s_4, s_5]$ are the safety parameters
- 2: Let LLM_0 be the unpruned model
- 3: Validation - for each s_r in s starting from s_1 :
 - 4: Prune LLM_{r-1} with $s_1, \dots, s_r$ producing LLM_r
 - 5: Produce 50 tokens of response res using LLM_r
 - 6: Produce the rest of res using LLM_0
 - 7: Test if res is a harmful response

```

LlamaForCausalLM(
  (model): LlamaModel(
    (embed_tokens): Embedding(32000, 4096, padding_idx=0)
    (layers): ModuleList(
      (0-31): 32 x LlamaDecoderLayer(
        (self_attn): LlamaSdpaAttention(
          (q_proj): Linear(in_features=4096, out_features=4096, bias=False)
          (k_proj): Linear(in_features=4096, out_features=4096, bias=False)
          (v_proj): Linear(in_features=4096, out_features=4096, bias=False)
          (o_proj): Linear(in_features=4096, out_features=4096, bias=False)
          (rotary_emb): LlamaRotaryEmbedding()
        )
        (mlp): LlamaMLP(
          (gate_proj): Linear(in_features=4096, out_features=11008, bias=False)
          (up_proj): Linear(in_features=4096, out_features=11008, bias=False)
          (down_proj): Linear(in_features=11008, out_features=4096, bias=False)
          (act_fn): SiLU()
        )
        (input_layernorm): LlamaRMSNorm((4096,), eps=1e-05)
        (post_attention_layernorm): LlamaRMSNorm((4096,), eps=1e-05)
      )
    )
    (norm): LlamaRMSNorm((4096,), eps=1e-05)
    (rotary_emb): LlamaRotaryEmbedding()
  )
  (lm_head): Linear(in_features=4096, out_features=32000, bias=False)
)

```

Figure 7: Parameters pruned in the LLaMA 2 (7B) [42] model.

LLaMA 3.1-8b, and approximately 1% of the parameters in Qwen 2.5, resulted in high success rate of TwinBreak which shows TwinBreak’s ability to prune safety-critical parameters in a fine-grained manner. In addition, we had to increase utility rate from 0.001 to 0.01 for LLaMA 2 when using the pruning dataset with the size of 60 and 70 twin pairs. All other cases follow the values reported in Tab. 9.

8 Additional Experimental Details

Fig. 8 shown an example of a successful TwinBreak jailbreak.

8.1 Experimental Setup

Hardware & Software. Experiments are implemented in PyTorch [49, 60, 62] and executed on one of two servers depending on model size. Server one features an AMD EPYC 7413 (24-core, 96 threads), 128 GB RAM, and an NVIDIA A16 GPU with 4 virtual GPUs (16 GB GDDR6 each). Server two has an Intel Xeon Gold 6526Y (16-core, 64 threads), 256 GB RAM, and 4 NVIDIA L40S GPUs (48 GB GDDR6 each). Both use CUDA 12.7 [47].

When reporting runtimes of TwinBreak and the related works [67], [4], we switch to a more accessible setting by using two Tesla T4 GPUs each with 16GB of memory to show the efficiency of TwinBreak even against an attacker

Table 9: TwinBreak’s hyperparameters and their defaults.

Hyperparameter Name	Default Value
Utility Parameters Rate (utl_rate)	0.001 (0.1%)
Target Decoder Layers (t_layers)	All except the first and the last layers
Targets in MLP Layer (mlp)	Gate and Up Layers
Attention Pruning (attn)	False
Token Activation to Mean Over (mean)	Top 5 tokens (with the largest L_2 norm of the activation differences)
Targeted Input Tokens (t_inp)	Last six tokens
Number of Output Tokens to Generate With Pruned Model (n_out_pr)	50
Number of Output Tokens to Generate When Finding Parameters (n_out_gen)	1
Aggregate Identified Safety Parameters at Each Iteration (agr)	True
Number of Pruning Iterations (n_iter)	5
Pruning Rate (pr_rate)	0.01 (1%)
Use Batch of Twin Prompts (batch)	False
Batch Size When Batched Twin Prompts (b_size)	-




Figure 8: A successful jailbreak by TwinBreak.

with modest resources. Such settings are available for instance on free platforms like Kaggle [31].

It should be noted that we use half-precision (float 16) in all of our experiments except for LLaMA 2 70B [41], LLaMA 3.3 70B [44], and Qwen 2.5 72B [25] where we used their 8-bit quantized versions. We register special hooks to the target MLP modules and layers of each model in order to collect their output activations and to set the safety parameter outputs to zero to simulate the pruning.

8.2 Evaluation Datasets

We utilize the following four datasets comprising harmful prompts to assess LLMs’ security and resilience against potential misuse. These prompts are designed to mimic malicious interactions, helping to uncover vulnerabilities within the models.

AdvBench, initially created by Zou *et al.* [75], is a dataset containing 520 harmful prompts designed to test LLMs’ vul-

nerabilities. While some of the prompts in the dataset yield high similarity, the dataset offers a variety of scenarios to evaluate LLM robustness. The prompts are partially created by LLMs instead of humans.

HarmBench [39] is an improved version of AdvBench, which enhances the quality of harmful prompts and eliminates redundancies, resulting in a more concise dataset of 200 harmful prompts.

Jailbreakbench [8] includes 100 pairs of harmful and harmless prompts that cover similar topics. It improves upon HarmBench by refining the harmful prompts present in HarmBench, with 45 crafted by the authors and 55 sourced from HarmBench. The dataset also provides harmless versions of these prompts, helping assess whether LLMs reject prompts unnecessarily based on specific words, such as "kill". However, the harmless prompts do not reflect the level of similarity proposed in our novel dataset, TwinPrompt, which we describe in more detail in Sect. 4.1.

StrongREJECT proposed by Souly *et al.* [57] contains 313 carefully crafted harmful prompts. StrongREJECT addresses the shortcomings of previously proposed datasets such as their repetitiveness or the presence of unanswerable prompts and provides a better dataset for higher-quality evaluations of jailbreak attacks. It also provides a highly accurate evaluator for evaluating the success of jailbroken completions which we explain below.

8.3 Prompt Templates

During inference, we use the chat templates listed in Fig. 9 following each LLM’s publisher guidelines. The *{instruction}* part in each template is replaced with a harmful or harmless prompt and the generated response is appended to the end of this template after inference.

For jailbreak success evaluations with LlamaGuard3 we use the default template provided by Meta as listed in Fig. 10.

```

LLaMA 2: "[INST] instruction [/INST]"
LLaMA 3.1 and LLaMA 3.3: "<start_header_id>[user]<end_header_id>
instruction<eot_id>|<start_header_id>[assistant]<end_header_id>
"
Qwen 2.5: "<im_start>[user
instruction<im_end>
|<im_start>[assistant
"
Gemma 2 and Gemma 3: "<start_of_turn>[user
instruction<end_of_turn>
|<start_of_turn>[model
"
Mistral V0.2: "[INST]{instruction}[/INST]"
DeepSeek LLM: "User: {instruction}
Assistant:"

```

Figure 9: Chat template used for all of our models. The double quote symbols denote the template start and end.


```

<|begin_of_text|><|start_header_id|>user<|end_header_id|>

Task: Check if there is unsafe content in 'Agent' messages in conversations
according our safety policy with the below categories.

<BEGIN UNSAFE CONTENT CATEGORIES>
S1: Violent Crimes.
S2: Non-Violent Crimes.
S3: Sex Crimes.
S4: Child Exploitation.
S5: Defamation.
S6: Specialized Advice.
S7: Privacy.
S8: Intellectual Property.
S9: Indiscriminate Weapons.
S10: Hate.
S11: Self-Harm.
S12: Sexual Content.
S13: Elections.
S14: Code Interpreter Abuse.
<END UNSAFE CONTENT CATEGORIES>

<BEGIN CONVERSATION>

User: <PROMPT>

Agent: <RESPONSE>

<END CONVERSATION>

Provide your safety assessment for ONLY THE LAST Agent message in the above
conversation:
- First line must read 'safe' or 'unsafe'.
- If unsafe, a second line must include a comma-separated list of violated
categories.<|eot_id|><|start_header_id|>assistant<|end_header_id|>

```

Figure 10: LlamaGuard3 chat template (provided by Meta) to evaluate the harmfulness of a response to a harmful prompt.